

Bhanu Prakash Bathini

New York, NY | (000) 000-0000 | your.email@example.com | [Portfolio](#) | [GitHub](#) | [Hugging Face](#)

EDUCATION

Northwest Missouri State University

Maryville, MO

Master of Science in Applied Computer Science

Aug. 2023 – Dec. 2024

- Coursework: Machine Learning, Data Mining, Database Systems, Cloud Computing, Software Engineering
- Certification: Claude Certified Architect (CCA) — Foundations, Anthropic (2026)

EXPERIENCE

Mondee

Fremont, CA

AI/ML Engineer

Jan. 2023 – Aug. 2026

- Built grounded document Q&A using Azure AI Search and GPT-4o with retrieval traces that cut hallucination escalations by 40%.
- Developed AI agents and automation workflows connecting models to FastAPI services and cloud endpoints for product teams.
- Improved groundedness scores by 35% by pinning source passages, refining prompts, and logging failure modes for review.
- Trained routing and classification models in PyTorch and scikit-learn that raised routing accuracy by 18% with sprint-level metrics.
- Prototyped AI flows from natural-language specs using Copilot and Cursor (~25% faster) and documented Azure AI setup for onboarding.

Cosmic Reality

Missouri, USA

AI/ML Engineer

Apr. 2020 – Oct. 2022

- Built PDF-to-structured-data pipelines in Python using PDFMiner and PyPDF2 to transform messy layouts into clean tables for models.
- Trained sentiment and spam classifiers with NLTK and scikit-learn, recovering ~12% accuracy after addressing class-imbalance issues.
- Developed classification and time-series models in PyTorch and TensorFlow with MLflow logging for replayable experiments.
- Packaged scoring applications with Docker on AWS EC2/S3 and documented deployment, shortening handoff time by about 40%.

PROJECTS

Agentic Workflow & Eval Harness | *Python, FastAPI, MLflow, n8n, Make.com, Retell AI*

2023 – 2026

- Built multi-step agents that call APIs, score outputs, and retry failed steps with MLflow logging for run comparisons.
- Prototyped workflow automations using n8n, Make.com, and Retell AI with prompt-engineering guardrails for reliable tool-to-human handoff.
- Developed evaluation harness tracking precision, recall, and groundedness metrics to systematically assess RAG and agent outputs.

ParseDoc.ai | *Python, PyTorch, Azure, Hugging Face, n8n*

2024 – 2026

- Extracted layout and text from long PDFs and trained detectors to produce structured fields for agents and search indexes.
- Generated labeled synthetic pages and improved bounding-box recovery by ~15%, publishing pipelines on Hugging Face and GitHub.
- Integrated extraction into n8n workflows and Azure Search to enable grounded document Q&A with auditable source traces.

RESEARCH, PUBLICATIONS & VOLUNTEER

Peer review: ACM RecSys; ISCA CATA (Springer) — reviews on method, evaluation, and clarity.

Research: Document layout detection and grounded Q&A (ParseDoc.ai); Hugging Face [dataset_parsecdoc.ai](#).

Volunteer: Leo Club charitable member — community service, fundraising, and local outreach.

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript

AI / ML: RAG, PyTorch, BERT fine-tuning, prompt engineering, TensorFlow, scikit-learn

Agents & Automation: n8n, Make.com, workflow orchestration, Retell AI, Claude, GitHub Copilot

Data & Integration: Pandas, APIs, FastAPI, document extraction, NumPy, data pipelines

Cloud & Dev Tools: Azure AI Search, Azure AI Foundry, Docker, GitHub Actions, AWS EC2/S3

Evaluation & Quality: MLflow, precision/recall, groundedness metrics, trace-based debugging