

Bhanu Prakash Bathini

Los Angeles, CA | (000) 000-0000 | your.email@example.com | [Portfolio](#) | [GitHub](#) | [Hugging Face](#)

EDUCATION

Northwest Missouri State University

Maryville, MO

Master of Science in Applied Computer Science

Aug. 2023 – Dec. 2024

- Coursework: Software Engineering, Database Systems, Cloud Computing, Data Mining, Machine Learning
- Certification: Claude Certified Architect (CCA) — Foundations, Anthropic (2026)

EXPERIENCE

Mondee

Fremont, CA

AI/ML Engineer

Jan. 2023 – Aug. 2026

- Built internal automation workflows connecting FastAPI backend services to cloud endpoints used by product and operations teams.
- Developed document Q&A API on Azure AI Search and GPT-4o with trace-based debugging that reduced hallucination escalations by 40%.
- Improved data quality scores by 35% through prompt refinement, source passage logging, and failure mode documentation for reviewers.
- Trained routing and classification models in PyTorch and scikit-learn; increased routing accuracy by 18% with precision and recall reporting.
- Used Copilot and Cursor to prototype backend services from specifications ~25% faster and documented workflows for team onboarding.

Cosmic Reality

Missouri, USA

AI/ML Engineer

Apr. 2020 – Oct. 2022

- Built Python data pipelines with PDFMiner and PyPDF2 that parsed thousands of malformed PDFs into structured tables for downstream systems.
- Developed classification models with scikit-learn and NLTK; improved accuracy by ~12% after addressing class imbalance and feature engineering.
- Trained PyTorch and TensorFlow models with MLflow experiment tracking to enable reproducible model evaluation and deployment workflows.
- Packaged services with Docker on AWS EC2/S3 and documented deployment steps, reducing handoff time by approximately 40%.

PROJECTS

Agentic Workflow & Eval Harness | *Python, FastAPI, MLflow, n8n, Make.com, Retell AI*

2023 – 2026

- Built multi-step backend agents with FastAPI that call APIs, evaluate outputs, and retry failed steps with logged metrics in MLflow.
- Prototyped workflow automation systems using n8n and Make.com with prompt-engineering guardrails for tool-to-human handoff visibility.
- Implemented evaluation harness tracking precision, recall, and failure cases to improve agent reliability across production workflows.

ParseDoc.ai | *Python, PyTorch, Azure, Hugging Face, n8n*

2024 – 2026

- Developed Python data pipeline extracting layout and text from PDFs so backend services receive structured fields instead of raw pages.
- Trained PyTorch detectors on synthetic labeled data and improved bounding-box recovery by ~15%; published datasets on Hugging Face and GitHub.
- Integrated extraction API with n8n workflows and Azure Search to enable grounded document Q&A with auditable source traces.

RESEARCH, PUBLICATIONS & VOLUNTEER

Peer review: ACM RecSys; ISCA CATA (Springer) — reviews on method, evaluation, and clarity.

Research: Document layout detection and grounded Q&A (ParseDoc.ai); Hugging Face `dataset_parsecdoc.ai`.

Volunteer: Leo Club charitable member — community service, fundraising, and local outreach.

TECHNICAL SKILLS

Languages: Python, JavaScript, SQL

Backend & Data: FastAPI, APIs, Pandas, NumPy, data pipelines, SQL joins

Cloud & DevOps: Docker, GitHub Actions, Azure AI Foundry, AWS EC2/S3

Automation & Tools: n8n, Make.com, workflow orchestration, GitHub Copilot, Cursor

AI / ML: PyTorch, scikit-learn, TensorFlow, MLflow, RAG